

Demo: Querying Cluster


Summary

Learn how to connect Query Editor v2 to an Amazon Redshift Cluster and run SQL queries on sample data.

Accessing AWS Redshift

Go to homepage and search for Redshift.


Click Serverless Dashboard in the sidepanel and open the Dashboard.



Click on the cluster name and then click Query Data → Query in query editor v2

Step 1 – Open Query Editor v2

- Opened **Query Editor v2** from the Redshift console.



- This is a web-based SQL editor used to query Redshift databases.

Step 2 – Connect to the Redshift Cluster


On the left side, the cluster was initially inactive.

Connection Method Used

- Temporary credentials using database username

Connection Details

Setting	Value
Database	dev
Username	awsuser



- A connection to the Redshift cluster was created.
 - The available databases became visible.
- | A Redshift Cluster is the actual database system where data is stored and queried.

Step 3 – View Available Databases

After connecting, the following databases appeared:

Native Databases

- dev
- sample_data_dev

External Databases

- awsdatalog
- dev@redshift-cluster-1

Inside the dev database:

- Default schema available:


- `public`

The `public` schema already contained sample Ticket tables:

- category
- date
- event
- listing
- sales
- users
- venue

These tables are used for practicing SQL queries in Redshift.

Step 4 – Use Notebook for Queries



- Opened the sample notebook from the left panel.
- Created a duplicate copy.


Notebook Name

`tickit-sample-notebook-for-cluster`

Purpose:

- Used to test SQL queries on the cluster database.
-

Step 5 – Error Encountered



Error Message :

schema tickit does not exist


Why did this happen?

The queries were trying to access tables using:

tickit.table_name

But the tables were actually stored inside the public schema.

Step 6 – Fix the Error



Removed:

- `search_path`
- `ticket.` prefixes from the queries

After this, queries directly used the tables from the `public` schema.

Step 7 – Run Validation Queries

Query 1 – Sales Per Event

Displayed:

- Event names
- Total sales amount collected for each event

Query 2 – Quantity Per Buyer

Displayed:

- Total quantity of tickets purchased by buyers
-

Final Result

- Queries executed successfully.
- Confirmed that:
 - The Redshift cluster connection works correctly
 - Sample Tickit data is available
 - Query Editor v2 can successfully query cluster databases